



Registration of Interest — AI Governance and Strategy

If you can't see an open role in the AI Governance & Strategy team that you want to apply for right now, complete this form to be notified of any roles that come along that meet your criteria. You'll receive an email when you register and then again once a job comes up that you might be interested in.

We tentatively plan to open a hiring round for multiple roles around September or October 2023 (though we may end up postponing that).

Prior to then, we may also highlight other potentially relevant opportunities to you (but you can opt out of that in the form or later). **We won't do anything else with your submitted form or provide other updates prior to the next time we're hiring.**

You can also fill in this [form](#) to register interest in working with RP's Existential Security team and/or fill in [this form](#) to register interest in roles on other

teams at RP.

Basic info about all our roles

- We invite anyone who is interested to apply, regardless of background, experience, or credentials. We will select candidates almost entirely based on their performance in our selection process, putting only minimal weight on CVs and references.
- Most of our positions are equally open to part- or full-time candidates (anywhere between 20-40 hours per week).
- The full-time annual salary we offer varies based on title and prior experience. The range is from \$62,000 (for some Research Fellows and Research Assistants) to \$118,500 USD (for some Senior Researchers or Senior Research Managers). Salaries for part-time staff are prorated for the fraction of 40 hours/week they work.
 - See [Titles at Rethink Priorities](#) for more information.
 - Payments may be made in different currencies and payment intervals depending on the location of applicants and legal requirements.

- We are a remote-first organization and welcome applicants from anywhere in the world.
 - We expect to be legally able to hire in most countries.
 - Most of our staff are in timezones between UTC-8 and UTC+3, but some staff are or have been in time zones outside that range and we have managed to make that work.
 - For Research Assistants and Research Managers (as opposed to Fellows and Researchers), we have a preference for candidates who are willing to take meetings — and sometimes respond to time-sensitive requests - during business hours for people in the UK, California, and Washington DC (UTC+0 to UTC+1 and UTC-8 to UTC-4). (But this doesn't necessarily require that a candidate be in a similar time zone themselves).
- Our Fellowships are 3-5 months, depending on each candidate's preference. All of our other roles are permanent.
 - Additionally, if a Fellow is potentially interested in a permanent position at Rethink Priorities, we are open to evaluating fellows during their time at RP for permanent roles. (Please email

info@rethinkpriorities.org if more information about this would make the difference between whether you apply or not.)

- For permanent staff, we offer comprehensive health coverage as regionally appropriate.
- Other benefits include stipends for workplace equipment, organization-wide mid- and end-of-year breaks, and generous leave policies (see “What we offer” below for details).

About Our Team

Our AIGS team tackles [a diverse set of questions](#) related to (1) what AI development and deployment scenarios may occur over the next few decades, and (2) how governments, firms, and other actors should prepare for, steer, and respond to various scenarios to reduce long-term/extreme risks. The ways we aim to positively impact the world include:

- informing major decisions by large grantmaking foundations and other important decision-makers interested in reducing long-term/extreme risks from AI

- identifying people who may excel in work related to AI governance (including roles in research, policy, grantmaking, etc.), helping them test their fit for such roles, and helping them build [career capital](#) for such roles

Currently, our team is organized into four main workstreams: China-West relations, compute governance, corporate labs, and US regulation & legislation.

We expect approximately half of our projects to be selected from a large set of project ideas that key decision-makers have directly told us they'd like us to do, while our other work will be more self-directed.

We expect most of our projects will result in at least one public output. However, we often deal with sensitive topics and projects, so some of our work will not be published, and we will expect our staff to be willing and able to handle confidential information with care.

You can see the current team members and read about them [here](#).

Department

AI Governance and Strategy

Employment Type

Flexible

Location

Remote

Workplace type

Fully remote

Compensation

\$62,000 - \$118,500 / year

Key Responsibilities

- Conduct research on important and diverse issues relevant to ensuring the development and deployment of AI is beneficial for society. Depending on the candidate, this could include activities such as:
 - Researching particular aspects of how AI development and deployment might proceed and what risks that might create (e.g., which AI technologies will be prone to rapid/large-scale proliferation and via what mechanisms)
 - Investigating how feasible and valuable various intervention options would be and how they could best be implemented (e.g., mechanisms for verifying AI-related treaties, or tools and

incentives for "whistleblowing" about risky development and deployment of AI systems)

- Conduct literature reviews, either to inform the team's work or also for sharing with external parties
- Find, read, summarize, and collate research, other relevant writings, and/or forecasting questions relevant to the team's work
- Writing accessible summaries on such topics tailored to particular audiences (e.g., US policymakers)
- Create concrete recommendations to advise the allocation of tens of millions of dollars a year among our interested audience of grantmakers, policymakers, and organization leaders.

Skills, Knowledge, and Expertise

- Problem-solving ability
- Clear and concise writing
- Attention to detail and a commitment to maintaining high quality and accuracy in all research output
- Ability to find, read, critically assess, and apply research from various disciplines and on various topics

- Ability to prioritize well and not drop the ball when given multiple tasks and deadlines (i.e., either hitting deadlines or proposing new deadlines well in advance)
- Comfort with independent, remote work (we're a 100% remote team across five different countries)

Useful but less important:

- Strong interpersonal skills and comfort with reaching out to various people outside of Rethink Priorities
- Data analysis skills
- Familiarity with G Suite (particularly [Documents](#) and [Sheets](#)), [Asana](#), and/or probabilistic modeling software (such as [Causal](#) or [Guesstimate](#))
- A disposition that favors evidence-based reasoning and critical thinking

We don't require any prior experience beyond the raw aptitude necessary to do well in our application process.

What we offer

- Salary between \$62,000 and \$118,500 pre-tax USD annually for a full-time position, prorated for

part-time work, as noted above.

- Flexible work schedule (what days and times you work is mostly up to you, especially for Researchers and Fellows)
- Flexible work location
- Generous benefits for permanent staff, including:
 - Comprehensive health coverage.
 - Stipends for workplace equipment
 - Organization-wide mid- and end-of-year breaks
 - Unlimited vacation and personal time off (several additional weeks in practice), generous family leave (six months paid at full salary, regardless of gender), and one “free week” per quarter for free exploration of potential research questions and/or projects that might be valuable to pursue
 - For temporary staff, up to 2 weeks of paid time off per 3 months at Rethink Priorities
 - 10% of your time can be devoted to professional development / learning / projects you want to pursue outside your assigned tasks
- Opportunities to grow/advance your career and engage in professional development, including specific opportunities to (further) test and build fit

for research roles, AI governance roles, and roles in the effective altruism community.

- We have a strong interest in providing staff with support on this, given that, as noted above, this is one of the main ways our team aims to positively impact the world. (That said, we're of course also excited to hire people who already have and know they have a strong fit for such roles.)
- Opportunity to contribute to a fast-growing, high-impact organization — our research is used by key decision-makers who influence the distribution of hundreds of millions of charitable dollars
- A caring team that values respectful work relations and a healthy work-life balance
- Low administrative bureaucracy
- We don't provide snacks but we could mail you a box of Oreos if you want!

About Rethink Priorities

Founded in 2018, [Rethink Priorities](https://rethinkpriorities.org/en/postings/95efd86b-2cc0-4cba-b394-660cdb4427f) (RP) is a nonprofit organization that addresses global priorities

—important and neglected issues—by researching solutions and strategies, mobilizing resources, and empowering our team and others. RP's mission is to generate the most significant possible impact for others in the present and the long-term future.

Our cause areas include animal welfare, global health and development, climate change, artificial intelligence, and other work to safeguard a flourishing long-term future. RP also aims to understand and support the professional communities working on these issues. Each researcher tends to focus on one particular cause area.

Rethink Priorities works as all of the following:

1. A consultancy doing commissioned work in response to demands from organizations doing high-impact work
2. A research institute driven by research agendas we set according to our own priorities.
3. A think tank aiming to inform public policy to improve the world.

4. An accelerator, incubator, and base for entrepreneurial projects.

Some of RP's recent accomplishments include:

- Publishing a nine-post sequence on [understanding the diffusion of large language models](#) which presents key findings from case studies on the diffusion of eight language models that are similar to GPT-3.
- Conducting and writing up results from [an expert survey on AI strategy](#), which has informed key decision-makers and been included in reading lists for people entering this field.
- Organizing a well-received summit for 35 leading members of the existential-risk-focused AI strategy and policy field.
- Producing public and nonpublic reports [on various topics](#), including [prospects for AI safety agreements between countries](#).
- Helping major foundations to answer their questions on climate change solutions, weather forecasting in lower- and middle-income countries, increasing access to medicine, and the effectiveness of [prizes](#) and other interventions.

- Comparing the capacity of [different animal species](#) to experience pleasure and pain to help philanthropists decide how to allocate funding.
- Investigating various [animal welfare interventions](#), as well as bringing to light the neglected areas of [invertebrate](#) and [insect](#) welfare.
- Publishing pieces on [nanotechnology](#) and [ways to use forecasting to improve the long-term future](#), as well as [supporting](#) those interested in these types of topics.
- Launching a [Special Projects Team](#) to incubate promising new initiatives, such as [Epoch](#) (a new AI research organization) and [Condor Camp](#) (longtermism movement-building in Brazil and Latin America).

We welcome you to review our database of published work [here](#).

We're supported by [Open Philanthropy](#), the [Survival and Flourishing Fund](#), and additional institutional and individual donors.

Additional information

If you'd like to know more about our hiring process, culture, and what working at Rethink Priorities is like, visit our [Career Opportunities](#) website.

Inclusivity and fairness: RP is committed to finding the best people for our team and to building an inclusive, equitable, and supportive community for you to thrive and do your best work. So please don't hesitate to apply for a role regardless of your age, gender identity/expression, political identity, personal preferences, physical abilities, veteran status, neurodiversity, or any other background. We provide reasonable accommodations and benefits, including flexible work schedules and locations, mental health coverage in medical benefits (as available), and technology budgets and professional development time that can be used to purchase assistive technology or engage in job coaching.

*We occasionally get asked about recommended reading materials. If you want, feel free to read more about our work [here](#). Note that members of our team specialize in different areas, and we hire for positions with varying levels of seniority, so please don't feel discouraged from applying if some of our work

seems too unfamiliar or advanced. If you are interested in further resources on longtermist research careers in general, Michael Aird, one of the AI Governance & Strategy team's Senior Research Managers, put together [this Forum post](#) with personal recommendations.

[Apply Now](#)

Our Hiring Process

Previous

[Next](#)

- [1](#)
- [2](#)
- [3](#)

[Find out more](#)

If there isn't a role available that aligns with your interests and expertise, we still want to hear from you. At a minimum, we'll keep your information on file, and let you know about future roles you might be a good fit for.

[Register Your Interest](#)[Apply Now](#)

Department

AI Governance and Strategy

Employment Type

Flexible

Location

Remote

Workplace type

Fully remote

Compensation

\$62,000 - \$118,500 / year

[View all opportunities at Rethink Priorities](#)[Apply Now](#)

[Rethink Priorities](#)  [Privacy Policy](#) [Cookies](#)
  
[Powered by Pinpoint](#)